

# On-Call Operational Runbook

Quick Reference Guide for Security-Layer Incidents  
Gateway, Risk-Engine, Mitigation, Audit, Notification

|                                  |                               |                                 |                            |
|----------------------------------|-------------------------------|---------------------------------|----------------------------|
| SEV-1 CRITICAL<br>5 min response | SEV-2 HIGH<br>15 min response | SEV-3 MEDIUM<br>60 min response | SEV-4 LOW<br>24 h response |
|----------------------------------|-------------------------------|---------------------------------|----------------------------|

## PURPOSE

This runbook provides immediate, actionable guidance for on-call engineers handling security-layer incidents. Keep this document accessible at all times during on-call shifts. Follow the standard checklist for every incident, then apply specific procedures based on incident type.

# 1. Grundprinzipien für On-Call

## 1.1 Ziele

- Sicherheitsvorfälle schnell erkennen
- Schaden minimieren
- Fraud verhindern
- System stabil halten
- Compliance sicherstellen

## 1.2 On-Call Verantwortlichkeiten

- Alerts entgegennehmen und bestätigen
- Diagnose sofort starten
- Sofortmaßnahmen einleiten
- Eskalieren wenn nötig
- Incident dokumentieren
- Post-Incident-Daten sammeln

## 1.3 Tools - Immer griffbereit

| Tool                 | Zugang                         | Zweck                            |
|----------------------|--------------------------------|----------------------------------|
| Grafana Dashboard    | grafana.cargobit.internal      | Metrics, Risk-Level Distribution |
| Alertmanager         | alertmanager.cargobit.internal | Active alerts, Silences          |
| Logs (Kibana/Loki)   | logs.cargobit.internal         | Gateway, Risk, Mitigation logs   |
| Deployment Dashboard | deploy.cargobit.internal       | Recent deployments, Rollback     |
| Runbooks             | runbooks.cargobit.internal     | Dieses Dokument + Details        |

Tabelle 1: On-Call Tools

## 2. Standard-Checkliste (für jeden Incident)

Diese Checkliste wird IMMER abgearbeitet, egal welcher Incident auftritt.

### 2.1 Sofort (0-2 Minuten)

| Schritt | Aktion                             | Check                                          |
|---------|------------------------------------|------------------------------------------------|
| 1       | Alert bestätigen                   | Alertmanager → Acknowledge                     |
| 2       | Severity prüfen                    | SEV-1 bis SEV-4                                |
| 3       | Betroffene Services identifizieren | Gateway, Risk, Mitigation, Audit, Notification |
| 4       | Traffic-Level prüfen (RPS)         | Baseline vergleichen                           |

Tabelle 2: Sofort-Checkliste

### 2.2 Diagnose (2-10 Minuten)

| Prüfung            | Tool             | Suche nach                        |
|--------------------|------------------|-----------------------------------|
| Logs               | Kibana / Loki    | ERROR, WARN, Exception            |
| Latenzen           | Grafana          | p99 > Baseline                    |
| Error-Rates        | Grafana          | > 0.1% = Warnung, > 1% = Kritisch |
| Queue-Lag          | Grafana          | > 2s = Warnung, > 5s = Kritisch   |
| Letzte Deployments | Deploy Dashboard | Änderungen in letzten 2h          |

Tabelle 3: Diagnose-Checkliste

### 2.3 Entscheidung (10-15 Minuten)

| Frage             | Ja →                          | Nein →                                 |
|-------------------|-------------------------------|----------------------------------------|
| Echter Incident?  | Weiter mit Maßnahmen          | Alert schließen, False Positive loggen |
| SEV-1 oder SEV-2? | Sofort eskalieren             | Normale Bearbeitung                    |
| Known Issue?      | Known-Issue-Playbook anwenden | Neue Diagnose starten                  |

Tabelle 4: Entscheidungs-Matrix

### 2.4 Maßnahmen (15-60 Minuten)

→ Siehe Incident-Typ-spezifische Abschnitte (3-7)

## 2.5 Recovery

- System stabilisieren (Metrics normalisieren)
- Alerts normalisieren (keine neuen Alerts)
- Audit-Trail prüfen (alle Events geschrieben)
- Queue-Lag < 2 Sekunden

## 2.6 Dokumentation

- Incident-Report ausfüllen (Template in Confluence)
- Timeline erstellen (von Alert bis Recovery)
- Logs sichern (für Post-Mortem)

## 3. Incident-Typ A: High-Risk Spike / Fraud-Welle

Severity: SEV-1 (Critical)

### 3.1 Symptome

- Viele RED-Decisions in kurzer Zeit
- Block-Rate steigt signifikant
- Notification-Service sendet viele High-Risk Alerts
- Support meldet ungewöhnlich viele Tickets

### 3.2 Diagnose

| Check                   | Grafana Panel             | Suche nach                  |
|-------------------------|---------------------------|-----------------------------|
| Risk-Level Distribution | Risk Dashboard → Level    | RED > 20% of total          |
| Triggered Rules         | Risk Dashboard → Rules    | Top 10 Rules identifizieren |
| Geo-Anomalien           | Risk Dashboard → Geo      | Regionale Häufung           |
| User/Company Cluster    | Risk Dashboard → Entities | Gleiche Company, IP-Range   |

### 3.3 Sofortmaßnahmen

| Maßnahme               | Kommando/Aktion                          | Wann                        |
|------------------------|------------------------------------------|-----------------------------|
| Strict Mode aktivieren | Gateway Config: strict_mode=true         | Bei bestätigter Fraud-Welle |
| Regel deaktivieren     | API: PUT /risk/rules/{id} → active=false | Bei feuender Regel          |
| Geo-Blockade           | Gateway Config: geo_block=["XX"]         | Bei regionalem Spike        |
| Benachrichtigung       | Slack: #incident-security                | Sofort bei SEV-1            |

### 3.4 Eskalation

- Security-Engineer: security-oncall@cargobit.eu
- Compliance: compliance@cargobit.eu

### 3.5 Recovery

- Strict Mode deaktivieren
- Risk-Rules neu trainieren/anpassen
- Fraud-Pattern in Regelwerk aufnehmen

## 4. Incident-Typ B: Risk-Engine Down / Degraded

Severity: SEV-1 (Critical)

### 4.1 Symptome

- Gateway Latenz steigt signifikant
- Viele RISK\_ENGINE\_UNAVAILABLE Errors in Logs
- Circuit-Breaker Status: OPEN
- Risk-Engine Healthcheck: ROT

### 4.2 Diagnose

| Check              | Tool/Kommando                   | Suche nach               |
|--------------------|---------------------------------|--------------------------|
| Health Check       | curl localhost:3003/risk/health | Status != 200            |
| Logs               | Kibana: service=risk-engine     | Error, Exception, OOM    |
| DB-Latenz          | Grafana: DB Dashboard           | p99 > 100ms              |
| CPU/Memory         | Kubernetes: kubectl top pods    | CPU > 80%, Mem > 90%     |
| Recent Deployments | Deploy Dashboard                | Änderungen in letzten 2h |

### 4.3 Sofortmaßnahmen

| Option               | Kommando/Aktion                                | Wann                  |
|----------------------|------------------------------------------------|-----------------------|
| Restart              | kubectl rollout restart deployment/risk-engine | Bei Service hängend   |
| Rollback             | kubectl rollout undo deployment/risk-engine    | Nach neuem Deployment |
| Fail-Safe aktivieren | Gateway Config: fail_safe=block                | Bei langem Ausfall    |

### 4.4 Eskalation

- Backend-Team Risk-Engine: backend-risk@cargobit.eu
- Platform-Team: platform-oncall@cargobit.eu

### 4.5 Recovery

- Circuit-Breaker schließen (Status: CLOSED)
- Latenzen normalisieren (p99 < 100ms)
- Audit-Trail prüfen (alle Events geschrieben)

## 5. Incident-Typ C: Mitigation-Queue Overload

Severity: SEV-2 (High)

### 5.1 Symptome

- Queue-Lag > 5 Sekunden
- Verzögerte 2FA-Verifikationen
- Verzögerte Delay-Mitigations
- Worker CPU hoch (> 80%)

### 5.2 Diagnose

| Check             | Grafana Panel                  | Threshold                 |
|-------------------|--------------------------------|---------------------------|
| Queue-Lag         | Mitigation Dashboard → Lag     | > 5s = Kritisch           |
| Worker Count      | Mitigation Dashboard → Workers | Vergleich mit Queue-Tiefe |
| Dead-Letter Queue | Mitigation Dashboard → DLQ     | > 100 = Cleanup nötig     |
| Mitigation Types  | Mitigation Dashboard → Types   | Anteil Delay/2FA/GPS      |

### 5.3 Sofortmaßnahmen

| Maßnahme           | Kommando/Aktion                                          | Priorität                   |
|--------------------|----------------------------------------------------------|-----------------------------|
| Worker skalieren   | kubectl scale deployment mitigation-worker --replicas=10 | 1 (Sofort)                  |
| Priorisierung      | Queue Config: priority=2FA>GPS>Delay                     | 2 (Wenn Lag > 10s)          |
| Delay deaktivieren | Mitigation Config: delay_enabled=false                   | 3 (Nur bei Extremsituation) |
| DLQ cleanup        | Admin Tool: dlq-flush                                    | 4 (Nach Stabilisierung)     |

### 5.4 Eskalation

- Mitigation-Service Team: mitigation-team@cargobit.eu
- Platform-Team: platform-oncall@cargobit.eu

### 5.5 Recovery

- Queue-Lag < 2 Sekunden
- Worker stabil (CPU < 70%)
- Delay-Mitigations wieder aktiviert

## 6. Incident-Typ D: Audit-Service Probleme

Severity: SEV-1 (Critical) - Compliance-relevant!

### 6.1 Symptome

- Audit-Write-Errors in Logs
- Audit-DB Latenz hoch
- Missing Audit Events Alerts
- Hash-Chain Validation Errors

### 6.2 Diagnose

| Check        | Tool                        | Suche nach     |
|--------------|-----------------------------|----------------|
| DB-Latenz    | Grafana: Audit DB Dashboard | p99 > 50ms     |
| Storage      | DB Admin: storage check     | > 80% full     |
| Write-Errors | Kibana: audit_write_error   | > 0 = Kritisch |

### 6.3 Sofortmaßnahmen

| Maßnahme          | Kommando/Aktion                                  | Hinweis                     |
|-------------------|--------------------------------------------------|-----------------------------|
| Restart           | kubectl rollout restart deployment/audit-service | Bei Service hängend         |
| DB-Failover       | DB Admin Console: Failover                       | Nur wenn Primary down       |
| Write-Queue flush | Admin Tool: audit-queue-flush                    | Gepufferte Events schreiben |

### 6.4 Eskalation

- Security-Engineer: security-oncall@cargobit.eu (Compliance-relevant!)
- DB-Team: db-team@cargobit.eu

### 6.5 Recovery

- Audit-Trail Validierung (alle Events vorhanden)
- Hash-Chain prüfen (Integrität bestätigt)
- Compliance benachrichtigen

## 7. Incident-Typ E: Notification-Service Failure

Severity: SEV-2 (High)



## 7.1 Symptome

- Slack Alerts kommen nicht an
- Email Delivery Failures
- Notification Queue wächst

## 7.2 Diagnose

| Check            | Tool                   | Suche nach       |
|------------------|------------------------|------------------|
| Delivery-Latency | Notification Dashboard | p99 > 30s        |
| Queue-Lag        | Notification Dashboard | > 60s = Kritisch |
| Provider-Status  | Provider Status Page   | Outage bekannt?  |

## 7.3 Sofortmaßnahmen

| Maßnahme          | Kommando/Aktion                                            | Hinweis                  |
|-------------------|------------------------------------------------------------|--------------------------|
| Provider wechseln | Notification Config: provider=fallback                     | Fallback-Slack/Email     |
| Worker skalieren  | kubectrl scale deployment notification-worker --replicas=5 | Bei Queue-Backlog        |
| DLQ prüfen        | Admin Tool: notification-dlq-check                         | Fehlgeschlagene Messages |

## 7.4 Eskalation

- Notification-Team: notification-team@cargobit.eu
- Platform-Team: platform-oncall@cargobit.eu

## 7.5 Recovery

- Delivery-Rate normal (> 99%)
- Queue-Lag < 2 Sekunden

## 8. Eskalationsmatrix

| Severity | Beschreibung                                    | Max. Reaktionszeit | Eskalation                       |
|----------|-------------------------------------------------|--------------------|----------------------------------|
| SEV-1    | Fraud-Welle, Risk-Engine Down, Audit-Failure    | 5 Minuten          | Security + Compliance + Platform |
| SEV-2    | Mitigation-Queue Overload, Notification Failure | 15 Minuten         | Security + Backend               |
| SEV-3    | Minor Errors, Slowdowns                         | 60 Minuten         | Backend                          |
| SEV-4    | Non-Critical                                    | 24 Stunden         | Team Lead                        |

Tabelle 5: Eskalationsmatrix

## 9. Kommunikationsplan

| Severity | Slack Channel      | Email an                       | Besonderheiten               |
|----------|--------------------|--------------------------------|------------------------------|
| SEV-1    | #incident-security | Security, Compliance, Platform | Incident Commander bestimmen |
| SEV-2    | #incident-backend  | Backend Team                   | Reguläre Updates             |
| SEV-3    | #dev-ops           | -                              | Ticket erstellen             |
| SEV-4    | -                  | -                              | Nur Ticket, kein Channel     |

Tabelle 6: Kommunikationsplan

## 10. Post-Incident Anforderungen

| Zeitraumen | Aufgabe                                | Verantwortlich      |
|------------|----------------------------------------|---------------------|
| 24h        | Timeline erstellen                     | On-Call Engineer    |
| 24h        | Root Cause identifizieren              | On-Call + Team      |
| 24h        | Impact Assessment                      | On-Call Engineer    |
| 48h        | Fix implementieren                     | Entsprechendes Team |
| 48h        | Regression Tests                       | QA Team             |
| 72h        | Post-Mortem Meeting                    | Team Lead           |
| 72h        | Lessons Learned dokumentieren          | On-Call Engineer    |
| 72h        | Risk-Rules aktualisieren (falls nötig) | Security-Engineer   |

Tabelle 7: Post-Incident Checkliste

## Key Contacts

| Team               | Email                          | Slack            |
|--------------------|--------------------------------|------------------|
| Security On-Call   | security-oncall@cargobit.eu    | #security-oncall |
| Backend On-Call    | backend-oncall@cargobit.eu     | #backend-oncall  |
| Platform On-Call   | platform-oncall@cargobit.eu    | #platform-oncall |
| Compliance         | compliance@cargobit.eu         | #compliance      |
| Incident Commander | incident-commander@cargobit.eu | #incidents       |

Tabelle 8: Key Contacts